

Modelling structured trial-by-trial variability in evidence accumulation

Steven Miletić, Niek Stevenson, Luke Strickland, Andrew Heathcote

22 June, 2026

In this tutorial, we will demonstrate how to design dynamic evidence accumulation models and fit them to data using hierarchical Bayesian methods using the *EMC2* package (Stevenson et al. 2026). Specifically, we will fit models to dataset 1 from Miletić et al. (2025), which corresponds to the ‘choice, short [CS]’ condition from Wagenmakers, Farrell, and Ratcliff (2004). In this experiment, participants were presented with a single digit and required to determine whether the digit was even or odd.

To enhance readability, the code blocks below include only the arguments strictly required to run the code. The R Markdown file that generates this document additionally contains hidden code blocks handling multi-core processing to speed up sampling and posterior predictive generation, automatic saving and loading of samples, and plot formatting to improve visibility in this PDF output. The source document is available at https://github.com/StevenM1/dynamic_tutorial/ (also linked on our OSF repository at <https://osf.io/4e3ta/files/github>), alongside an HTML version that facilitates copy-and-pasting.

First, let’s load the package and data:

```
# remotes::install_github("ampl-psych/EMC2@trend_api", dependencies=TRUE, Ncpus=8,
# figure.args=c('--native')); .rs.restartR()
library(EMC2)
set.seed(1) # for reproducibility

# This is some code that contains plotting functions used later on in this tutorial
code_url <- paste0(
  "https://raw.githubusercontent.com/StevenM1/",
  "dynamic_tutorial/refs/heads/main/plotting_utils.R")
source(code_url)

# Load in the data
data_url <- paste0(
  "https://github.com/StevenM1/dynamic_tutorial/",
  "raw/refs/heads/main/datasets/dataset_1.RData")
load(url(data_url))

# Inspect the data
print(head(dat))
```

```
#>  subjects    S    R    rt
#> 1         1 even even 0.542
#> 2         1 even even 0.426
#> 3         1 even even 0.501
#> 4         1 even even 0.403
#> 5         1  odd  odd 0.497
#> 6         1 even even 0.518
```

Here, **subjects** refers to subject number, **S** to the presented stimulus (even/odd), **R** the response (even/odd), and **rt** the response time in seconds.

0.1 Baseline models

We begin by setting up two static baseline models. One is a racing diffusion model (RDM) and the other a Wiener diffusion model (WDM), which refers to the diffusion decision model (DDM) without between-trial variability (sometimes called the ‘simple’ DDM). Using these models, we demonstrate a minimal but principled analysis pipeline (see Figure 1): model specification, prior specification, model estimation (fitting), model criticism (checking convergence and quality of fit), inference (in examples below, we focus on model comparison), and validation (here, a simple parameter recovery). For more complex examples, we refer to Strickland et al. (2026). We then show how the same pipeline applies when model specifications are extended to include time-varying parameters.




Figure 1: Overview of the modelling workflow in *EMC2*. Solid boxes indicate modelling steps, with key functions listed above or beside the box. Functions relating to time-varying parameters are highlighted in green. Note that not all functions will be covered in this tutorial. The `DDM()`, `RDM()`, `LBA()`, and `LNR()` refer to the choice models that are supported in *EMC2* at the time of writing.

In *EMC2*, model specification is done via the `design()` function. It minimally requires: 1) the relevant experimental design factors (most easily specified by just passing the data to which the design will be applied, as we do below), 2) the model (e.g., `RDM`, `DDM`), and 3) a formula (in *R*’s linear modelling language) specifying how model parameters vary with design factors. The present dataset has one design factor **S** (for stimulus) with two levels: odd and even. *EMC2* codes accumulators (in race models) and bounds (in the `WDM/DDM`) with the levels of the response **R**. In our example, we used response coding (which we recommend in general), so the drift rate should vary with **S** — the evidence for an ‘odd’ response differs between odd and even trials. We can enforce the `WDM` drift rate to be equal in magnitude but opposite in sign across the two stimulus levels by specifying a contrast matrix. In a contrast matrix, column names represent sampled parameters and values indicate how those parameters are linearly combined to obtain the model parameter in each design cell (rows). Here we define `Smat` with two rows, one per factor level:

```
# set up contrast matrix that flips the drift rate between the stimulus levels
Smat <- matrix(c(-1, 1), nrow=2, dimnames=list(NULL, "dif"))
design_WDM <- design(model=DDM, # Pass model type
                    data=dat,   # Pass data
                    # the contrast matrix should affect factor S
                    contrasts=list(S=Smat),
                    formula=list(Z~1, v~S, a~1, t0~1))
```

```
#>
#> Sampled Parameters:
#> [1] "Z"      "v"      "v_Sdif" "a"      "t0"
#>
#> Design Matrices:
#> $Z
#> Z
#> 1
#>
#> $v
#> S v v_Sdif
#> even 1 -1
#> odd 1 1
#>
#> $a
#> a
#> 1
#>
#> $t0
#> t0
#> 1
#>
#> $s
#> s
#> 1
#>
#> $st0
#> st0
#> 1
#>
#> $sv
#> sv
#> 1
#>
#> $SZ
#> SZ
#> 1
```

```
# The documentation in ?DDM lists all the parameters and their estimation scales
```

This process of combining *sampled* parameters into *estimated* parameters is what we call *design mapping* — this will play an important role later when specifying trends. By default, `design()` prints the design matrices (which we will turn off after the next example by setting `report_p_vector=FALSE` to limit the output length), but `mapped_pars()` provides some additional information:

```
mapped_pars(design_WDM)
```

```
#> $v
#> S
#> even : v - v_Sdif
#> odd  : v + v_Sdif
```

Note that, following *R*'s linear modelling language, an intercept is automatically added for drift rate *v*. Here, the *sampled* parameter *v* corresponds to an overall drift bias towards the upper bound. We can choose to set this as a constant to 0 by adding a constants argument to `design()`, as we will do in later examples.

For race models like the RDM, *EMC2* creates a data-augmented design matrix (`dadm`) under the hood - a data frame with one row *per accumulator* per trial per subject. `dadms` always have a factor column `lR` indicating which *response* each accumulator codes for ('odd' or 'even' here). When a `matchfun` is provided, a second factor column `lM` is added to indicate whether the accumulator matches the correct response. This enables a mean-difference parameterisation of the two drift rates: one parameter for mean evidence and one for the difference. We implement this via a contrast matrix `ADmat` with values -0.5 and 0.5 for the mismatching and matching accumulator rows, respectively:

```
# set up contrast matrix for mean-difference parametrisation
ADmat <- matrix(c(-.5, .5), ncol=1, dimnames=list(NULL, "d"))

design_RDM <- design(model=RDM,
                    data=dat,
                    contrasts=list(lM=ADmat), # the ADmat concerns factor lM
                    # matchfun tells the design which accumulators correspond to the correct
                    # answer
                    matchfun=function(d) d$S==d$lR,
                    formula=list(B~1, v~lM, t0~1))
```

```
#>
#> Sampled Parameters:
#> [1] "B"      "v"      "v_lMd" "t0"
#>
#> Design Matrices:
#> $B
#> B
#> 1
#>
#> $v
#>      lM v v_lMd
#> TRUE 1  0.5
#> FALSE 1 -0.5
#>
#> $t0
#> t0
#> 1
#>
#> $A
#> A
#> 1
#>
#> $s
#> s
#> 1
```

Note that the design matrix for drift rates *v* is similar to the WDM case, but the rows now correspond to the matching and mismatching accumulator rather than the two stimulus levels. Inspecting `mapped_pars()`:

```
mapped_pars(design_RDM)
```

```
#> $v
#> 1M
#> TRUE : exp(v + 0.5 * v_1Md)
#> FALSE : exp(v - 0.5 * v_1Md)
```

The output again looks similar, but reveals something perhaps unexpected: the linear combination of v and v_{1Md} is exponentiated. This is default *EMC2* behavior: parameters that must be strictly positive (such as RDM drift rates, and thresholds and non-decision times in both models) are estimated on the log scale, since negative values either have no known analytical likelihood (RDM drift rates) or a likelihood of 0 (thresholds and non-decision times). Additionally, some parameters (such as the DDM's start point relative to its bound) are estimated on the probit (inverse cumulative normal) scale to enforce that they lie between 0 and 1 after transformation. Default transformations are documented in `?DDM`, `?RDM`, `?LBA`, and `?LNR`, and can be overridden via the `transform` argument of `design()`. The key point is that sampled parameters are first *mapped* to design cells and then *transformed* to yield the final parameter value used to compute each observation's likelihood — a distinction that has implications for specifying time-varying parameters, as demonstrated later.

EMC2 uses a particle-within-Gibbs sampler that assumes a multivariate Gaussian group-level distribution with conjugate priors to facilitate Gibbs sampling. This requires priors on both the group-level mean vector and the covariance matrix. For the group-level means, *EMC2* places independent Gaussian priors with user-specified means and variances, defaulting to $N(0, 1)$ on the sampled scale. One benefit of the transforms is that the $N(0, 1)$ default prior always has support only over valid parameter values on the natural scale. The `plot_prior()` function visualises the corresponding implied prior on the natural scale:

```
# prior() creates a prior object from a design; defaults to N(0,1) on the sampled scale
prior_RDM <- prior(design_RDM)
# plot() visualises the implied prior on the natural scale
plot(prior_RDM)
```




Figure 2: Implied default priors for the group-level mean parameters of the baseline RDM. v_{1MTRUE} and $v_{1MFALSE}$ refer to the drift rate of the accumulators matching ('correct') and mismatching ('error') the stimulus, respectively; B to the threshold, and t_0 to the non-decision time.

For the covariance matrix, *EMC2* offers several options; by default, the full covariance structure is estimated using a scale-mixture inverse-Wishart prior with inverse-Gamma mixture weights (Huang and Wand 2013), which induces weakly informed half- t priors on the standard deviations and uniform priors on the correlations. In the examples below, we use the default prior settings for the covariance matrix throughout; for a full

treatment of the available options, see Stevenson et al. (2026) Priors can be overridden with the `prior()` function, as documented in `?prior` and demonstrated for group-level means in examples below.

Next, we can combine the data, design, and priors, into an `emc` object using the `make_emc` function, and pass the `emc` object to the fit function for sampling (by default `make\_emc` sets up sampling of 3 chains):

```
# make_emc combines the data, design, and optionally custom priors into a single object, which
can then be passed to fit() to start sampling
emc_rdm <- make_emc(dat, design_RDM, prior_list=prior_RDM)
emc_wdm <- make_emc(dat, design_WDM) # when not specified, the default priors are used

emc_rdm <- fit(emc_rdm)
emc_wdm <- fit(emc_wdm)
```

After fitting, it is important to check convergence. The `check()` function plots the group-level mean, group-level variance, and subject-level parameters with the highest R-hat (i.e., the worst cases), and it prints convergence statistics to the console: Rhat (values near 1 indicate the chains are well mixed), and ESS (effective sample size), which takes account of redundancy among the samples. Convergence is supported and even in the worst case redundancy among the posterior parameters (by default 1000 per chain) is relatively mild (thanks to *EMC2s* adapt stage, which increases sampling efficiency).

```
# check() plots trace plots for the worst-converging parameters and prints convergence
statistics (Rhat, ESS) to the console
check(emc_rdm)
```

```
#> Iterations:
#>      preburn burn adapt sample
#> [1,]      0    0    0   1000
#> [2,]      0    0    0   1000
#> [3,]      0    0    0   1000
#>
#>  mu
#>      B      v      v_lMd      t0
#> Rhat  1.001  1.001  1.002  1.005
#> ESS  3191.000 2342.000 2552.000 2738.000

#>
#>  sigma2
#>      B      v      v_lMd      t0
#> Rhat  1.002  1      1      1.002
#> ESS  1910.000 1347  1352 1397.000

#>
#>  alpha highest Rhat : 4
#>      B      v      v_lMd      t0
#> Rhat  1.002  1.003  1.003  1.004
#> ESS  2035.000 1842.000 1703.000 2156.000
```




Figure 3: Trace plots of the sampled parameters with highest Rhat values (i.e., worst convergence) at the group-level mean (left, μ), the group-level variance (middle, σ^2), and the subject-level (α ; in this case, subject number 4). ESS = effective sample size.

Superimposed prior and posterior distributions and credible intervals can be obtained with `plot_pars()` and `credint()`, respectively. Both accept a selection argument specifying which type of parameters to use — by default “mu” (group-level means), with alternatives “alpha” (subject-level), “sigma2” (group-level variance), and “correlation” (group-level correlations).

```
# plot_pars() plots posterior and prior densities, by default selection="mu" (group-level means)
plot_pars(emc_rdm)
```




Figure 4: Posterior density (black) and prior density (red) of sampled group-level mean parameters of the baseline RDM. Contraction indicates the reduction of uncertainty in the posterior versus the prior (1 minus the ratio of the posterior to prior variance).

By default, the `credint()` function returns the credible intervals for the sampled parameters, but it can also return the credible intervals of implied posteriors after mapping and transforming:

```
# credint() returns credible intervals for the sampled parameters by default
credint(emc_rdm)
```

```
#> $mu
#>      2.5%    50%   97.5%
#> B      0.006  0.294  0.570
#> v      0.320  0.669  0.999
```

```
#> v_lMd 0.946 1.562 2.205
#> t0 -1.977 -1.668 -1.388
```

```
# map=TRUE returns intervals on the natural scale (after mapping and transforming)
credint(emc_rdm, map=TRUE)
```

```
#> $mu
#>          2.5%   50% 97.5%
#> v_lMFALSE 0.909 1.069 1.237
#> v_lMTRUE  4.205 4.323 4.446
#> B         1.339 1.405 1.474
#> t0        0.186 0.195 0.203
```

Supplying the fit to the `predict()` function generates posterior predictives, which can then be used to plot the global quality of fit with defective cumulative density function (CDF) plots:

```
# the predict() function takes a fitted model and generates (by default) 50 replicates of the
# full dataset with 50 randomly chosen draws from the posterior
pp_rdm <- predict(emc_rdm)
pp_wdm <- predict(emc_wdm)

par(mfrow=c(1, 4))
plot_cdf(emc_rdm, pp_rdm,
  # we request a separate CDF for each level of the factor S (stimulus)
  factors=c("S"),
  # EMC2's plot_cdf() by default sets a reasonable par();
  # here, we override this to plot the two models side-by-side
  set.par=FALSE)
plot_cdf(emc_wdm, pp_wdm, factors=c("S"), set.par=FALSE)
```




Figure 5: Defective cumulative density plots of the baseline RDM (left two columns) and WDM (right two columns). Circles indicate the 10th, 50th, and 90th quantiles of the response time distributions for each response type R (solid = 'even', dashed = 'odd') in both conditions of the stimulus S (panels) for the data (black) and posterior predictives (green). The cumulative density functions are made defective as they sum to the proportion of each choice (rather than 1).

The RDM shows a slightly better fit, most visibly in the accuracy and 90th RT quantile. Formal model comparisons using DIC, BPIC, or marginal deviances (note that preferred models have lower DIC (Spiegelhalter et al. 2002), BPIC (Ando 2007) or MD ($= -2 * \text{marginal likelihood}$) values) can be obtained with `compare()`:

```
compare(list(rdm=emc_rdm, wdm=emc_wdm))
```

```
#>      MD wMD   DIC wDIC  BPIC wBPIC EffectiveN meanD Dmean minD
#> rdm -6577   1 -6680    1 -6651     1         29 -6708 -6732 -6737
#> wdm -5258   0 -5369    0 -5337     0         32 -5402 -5422 -5434
```

The posterior predictives can also be used for a parameter recovery study. By default, `predict()` generates 50 datasets (each tagged with a unique index `postn`) along with the data-generating parameters as an attribute. We can extract any generated dataset with its corresponding parameters, refit the same model, and assess whether the estimated parameters recover the data-generating ones:

```
# select a single simulated dataset and its data-generating parameters
postn <- 1
simulated_data <- pp_rdm[pp_rdm$postn==postn,]
true_pars <- attr(pp_rdm, "pars")[[postn]]

# combine them again in an EMC object, and fit
emc_simulated_rdm <- make_emc(simulated_data, design_RDM)
emc_simulated_rdm <- fit(emc_simulated_rdm)

# The quality of recovery can be plotted with recovery(). Note that we test for quality of
# subject-level parameter recovery here, so we select selection="alpha"
recovery(emc_simulated_rdm, true_pars, selection="alpha")
```




Figure 6: Parameter recovery results for the subject-level parameters. Estimated parameters (y-axis) are plotted against data-generating (x-axis). Error bars indicate 95% credible intervals, circles are on the median posterior, and diagonal lines is the identity. Pearson correlation between data-generating and posterior median. RMSE = root mean squared error.

The code above is purely meant for demonstrative purposes; here, we only simulated six participants following the data set size, but real parameter recovery studies should rely on more to cover a wider range of parameter settings. *EMC2* Also has built-in routines for simulation-based calibration that can be run on any models - for this functionality, see `?run_sbc`.

Together, these steps illustrate the key elements of a minimal but principled modelling workflow. In the following sections, we turn to time-varying parameters. The same functions apply throughout — the main difference lies in model specification, with some additional considerations for prior specification that we highlight along the way.

0.2 Descriptive trends

Both the descriptive trends and the formal mechanisms of dynamics work through **trend** objects in *EMC2*. A trend object combines two building blocks: a **kernel** and a **base**. The kernel specifies (1) the covariate of interest (e.g., time on task) and (2) the function applied to it (e.g., linear, exponential, power, polynomial, or delta rule). The base specifies (3) which decision parameter is informed by the kernel’s output, and (4) the functional form of that mapping. The `trend_help()` function gives an overview of the available options: The `trend_help()` function gives an overview of the options:

```
trend_help()
```

```
#> Available kernels:
#>   custom: Custom C++ kernel: provided via register_trend().
#>   lin_decr: Decreasing linear kernel: k = -c
#>   lin_incr: Increasing linear kernel: k = c
#>   exp_decr: Decreasing exponential kernel: k = exp(-d_ed * c)
#>   exp_incr: Increasing exponential kernel: k = 1 - exp(-d_ei * c)
#>   pow_decr: Decreasing power kernel: k = (1 + c)^(-d_pd)
#>   pow_incr: Increasing power kernel: k = 1 - (1 + c)^(-d_pi)
#>   poly2: Quadratic polynomial: k = d1 * c + d2 * c^2
#>   poly3: Cubic polynomial: k = d1 * c + d2 * c^2 + d3 * c^3
#>   poly4: Quartic polynomial: k = d1 * c + d2 * c^2 + d3 * c^3 + d4 * c^4
#>   delta: Standard delta rule kernel: k = q[i].
#>           Updates q[i] = q[i-1] + alpha * (c[i-1] - q[i-1]).
#>           Parameters: q0 (initial value), alpha (learning rate).
#>   delta2lr: Dual learning rate delta rule: k = q[i].
#>           Like the standard delta rule, but with separate
#>           learning rates for positive and negative prediction errors.
#>           Parameters: q0 (initial value), alphaPos (learning rate for positive PEs),
#>           alphaNeg (learning rate for negative PEs).
#>
#> Available base types:
#>   lin: Linear base: parameter + w * k
#>   centered: Centered mapping: parameter + w*(k - 0.5)
#>   add: Additive base: parameter + k
#>   identity: Identity base: k
#>
#> Phase options:
#>   premap: Trend is applied before parameter mapping. This means the trend parameters
#>           are mapped first, then used to transform cognitive model parameters before
#>           their mapping.
#>   pretransform: Trend is applied after parameter mapping but before transformations.
#>           Cognitive model parameters are mapped first, then trend is applied,
#>           followed by transformations.
#>   posttransform: Trend is applied after both mapping and transformations.
#>           Cognitive model parameters are mapped and transformed first,
#>           then trend is applied.
```

Linear, exponential, and power kernels are provided in both increasing and decreasing versions (e.g., `exp_incr` and `exp_decr`). In simple applications, the direction of a trend can also be flipped by allowing the associated weight parameter to be positive or negative, so the choice between “incr” and “decr” is not strictly necessary. However, having separate increasing and decreasing kernels is useful when one wants to enforce a direction of effect a priori. For example, one might hypothesise that a practice effect must increase a parameter over trials (e.g., a growing drift-rate difference) and wish to rule out decreasing trends by construction. In Section 0.2.2.1 we illustrate how this can be achieved by combining directional kernels with suitable transformations on the weight parameters (e.g., sampling the weight on the log scale to restrict it to be non-negative).

The `phase` argument controls *when* the trend is applied in the mapping pipeline, and has three options: "`premap`" (default), "`pretransform`", and "`posttransform`". To understand why this matters, recall that *EMC2* uses model formula language to create a design matrix that maps sampled parameters to each design cell, translating them into the "core" parameters per accumulator per trial. In the case of the RDM, these core parameters are `v`, `B`, `t0`, `s`, and `A` (the omission of `A` in the design above fixes it to zero). Crucially, many parameters of interest, such as between-accumulator differences (e.g., `v_1Md` in the examples above) or between-condition differences (e.g., the effect of speed-accuracy trade-off cues on thresholds), only exist before this mapping step. For such parameters, the trend must be applied prior to mapping, which is why "`premap`" is the default.

However, this has an important implication: by default, *EMC2* estimates parameters that must be strictly positive (e.g., thresholds, non-decision times, RDM drift rates) on the log scale, and parameters bounded between 0 and 1 on the probit scale. Since mapping occurs before transformation, a trend applied "`premap`" will subsequently be passed through these nonlinear transforms — meaning a linear trend on the sampled scale becomes, for example, an exponential trend on the natural scale (as illustrated below).

There are two ways to obtain a linear trend on the natural scale. First, setting `phase="posttransform"` in the base causes the trend to be applied after both mapping and transformation, so the covariate acts directly on the natural-scale parameter. Note, however, that "`posttransform`" cannot be used for parameters that only exist before mapping, such as `v_1Md`. Second, the log transform for the target parameter can be disabled entirely via the `transform` argument of `design()`, so that the parameter is sampled on its natural scale; the trend can then be applied "`premap`" without any subsequent nonlinear transformation. Both approaches are illustrated below.




Figure 7: Illustration of parameter mapping and transformation in EMC2. The estimated parameter vector is first mapped, then transformed. In the illustration, only drift rates `v` and thresholds `B` are shown. In this example, the trend is applied `premap`, which leads the threshold to increase with trials (`t`). After mapping, the parameters are expanded to each design cell, with in this example two accumulators per trial (one for each latent response `1R`, and one matches the stimulus `1M=T`). Finally, the parameters are transformed to the natural scale. Since in this example, a linear trend is applied `premap`, this leads to an exponentially increasing trend `posttransform`.

0.2.1 Linear trends

We illustrate this with a concrete example. To impose a linear trend on thresholds in the RDM, we first add trial number as a covariate:

```
# Add a 'trials' column specifying trial number
dat <- EMC2::add_trials(dat)

# Rescale the effect of this covariate to a larger parameter range to help sampling and
# interpretation (there are 1024 trials so the corresponding estimated parameter (B.w, see below)
# indexes the change from start to finish of the experiment).
dat$trials2 <- dat$trials/1024
```

Which we can then specify as a covariate in `make_kernel()`:

```
lin_kernel <- make_kernel(type="lin_incr",      # kernel type
                          cov_names="trials2") # covariate of interest
lin_base <- make_base(type="lin",             # base type
                      target_parameter="B",    # target parameter
                      kernel=lin_kernel,
                      phase="premap")          # and phase
lin_trend <- make_trend(lin_base)

RDM_lin_B <- design(model=RDM,
                    data=dat,
                    contrasts=list(lm=ADmat),
                    matchfun=function(d) d$S==d$1R,
                    formula=list(B~1, v~1M, t0~1),
                    covariates="trials2",      # specify relevant covariate columns
                    trend=lin_trend,           # add trend
                    report_p_vector=FALSE)    # suppress outputs for brevity
```

Note that we are now sampling one extra parameter compared to before, *B.w*, which is the weight of the influence of rescaled trials on thresholds: $B_t = B_0 + B.w * trials2$. Plotting the resulting trial-by-trial threshold with `plot_trend()` (Figure 8) confirms the expected exponential shape — a direct consequence of the "premap" phase combined with the default exp transform on B, as described above.

```
emc_lin <- make_emc(dat, design=RDM_lin_B)
p_vector <- c(B=1, v=1, v_1Md=1,
              t0=0.1, B.w=1)

# plot_trend() visualises the trial-by-trial parameter values implied by a given parameter
# vector. par_name specifies which parameter to plot; filter selects a subset of the dadm rows
# (here: the "odd" accumulator only as otherwise the identical trend for the "even"
# accumulator is also plotted).
plot_trend(p_vector, emc=emc_lin,
            par_name="B", subject=1,
            filter=function(data) data$1R=="odd",
            main="Threshold for odd")
```




Figure 8: Trialwise threshold parameters when using a linear effect of trials at the premap stage. The subsequent default exponentiation leads to a non-linear effect of trial number on thresholds.

`mapped_pars()` makes this explicit:

```
mapped_pars(RDM_lin_B)
```

```
#> $B
#>
#> intercept : exp(B_t)
#> Trends:
#> B_t = B + w * trials2
#>
#> $v
#> lm
#> TRUE : exp(v + 0.5 * v_lMd)
#> FALSE : exp(v - 0.5 * v_lMd)
```

The two solutions described above (`phase="posttransform"` and disabling the exp transform via `transform=list(func=c(B="identity"))`) both recover the intended linear trend on the natural scale, as shown below:

```

# Sample on log scale, but apply trend posttransform
lin_base_post <- make_base(type="lin", target_parameter="B",
                           kernel=lin_kernel,
                           phase="posttransform") # change phase
lin_trend2 <- make_trend(lin_base_post)
RDM_lin_B2 <- design(model=RDM,
                     data=dat,
                     contrasts=list(lM=ADmat),
                     covariates="trials2",
                     matchfun=function(d) d$S==d$lR,
                     formula=list(B~1, v~lM, t0~1),
                     trend=lin_trend2,
                     report_p_vector=FALSE)

# Alternatively, sample on natural scale, and apply trend premap
RDM_lin_B3 <- design(model=RDM,
                     data=dat,
                     contrasts=list(lM=ADmat),
                     covariates="trials2",
                     matchfun=function(d) d$S==d$lR,
                     # ...and tell EMC2 to sample the threshold on the natural scale
                     transform=list(func=c(B="identity")),
                     formula=list(B~1, v~lM, t0~1),
                     trend=lin_trend, # original trend
                     report_p_vector=FALSE)

# Plot both trends with the same parameter vector
emc_lin2 <- make_emc(dat, design=RDM_lin_B2)
emc_lin3 <- make_emc(dat, design=RDM_lin_B3)
p_vector <- c(B=1, v=1, v_lMd=1, t0=0.1, B.w=1)

par(mfrow=c(1, 2))
plot_trend(p_vector, emc=emc_lin2, par_name="B", subject=1,
           filter=function(data) data$lR=="odd",
           main="Posttransform")
plot_trend(p_vector, emc=emc_lin3, par_name="B", subject=1,
           filter=function(data) data$lR=="odd",
           main="Premap, identity")

```




Figure 9: Trialwise threshold parameters when using a linear effect of trials at the posttransform stage (left), or at the premap stage without an exponentiation (right).

Both yield the expected linear effect (but note the differences on the y-axis, arising due to the presence of the exponentiation in the `posttransform` case).

As a cautionary note, keep in mind that the default priors in *EMC2* are Gaussian distributions centred on 0 with unit variance. Since many cognitive model parameters cannot be negative (and by default *EMC2* will still enforce per-parameter bounds, see help for the `bound` argument in `?design`), a $N(0,1)$ prior is poorly chosen for those parameters that do not have support on the real line. For estimation and sampling, this usually has little influence in practice, but it may be important when estimating Bayes factors, since they marginalise the likelihood over the prior distribution, which now includes impossible negative values.

Based on these considerations, we offer the following practical recommendations. When the target parameter only exists before mapping (e.g., between-accumulator or between-condition difference parameters such as `v_lMd`), `"premap"` is the only option. If the trend is also required to have a particular form on the natural scale that is incompatible with the default transform (for instance, because the hypothesis under test implies a linear relationship) we recommend disabling the default transformation via `transform=list(func=c(B="identity"))` (where `B` should be replaced with the target parameter). Users should then take care to specify priors that are appropriate for the natural scale of the parameter, and be aware that Bayes factors will marginalise over a prior that now possibly includes values in the impossible negative range.

When the target parameter exists after mapping, `"posttransform"` is a simpler and safer choice, as the trend acts directly on the natural-scale parameter. In this case, keeping or disabling the default transformation is largely a matter of preference: the log or probit scaling can in principle aid sampling by keeping the parameter space unbounded, but in our experience the practical difference is small. Whether or not transformations are used, obtaining meaningful Bayes factors for trend parameters requires carefully specified, informed priors — a requirement that is arguably not specific to trend models, but applies broadly in Bayesian cognitive modelling.

With all that in mind, we can now start sampling our first model, `RDM_lin_B3`. Priors were left at their default values with one exception in light of the issues just discussed, a more informative ($N(2,0.5)$) prior on the threshold `B`, with an increased mean consistent with our experience with the RDM, and a decreased standard deviation to further reduce the prior density on negative thresholds, as shown in Figure 10.

```
# mu_mean and mu_sd override the default N(0,1) for the specified parameters
prior_linB3 <- prior(RDM_lin_B3, mu_mean=c(B=2), mu_sd=c(B=0.5))
emc <- make_emc(dat, design=RDM_lin_B3, prior_list=prior_linB3)

# map=FALSE plots the prior on the sampled scale rather than the natural scale
plot(prior_linB3, map=FALSE)
```




Figure 10: Prior distributions for the group-level mean parameters of the RDM with a linear trend, with the threshold `B` sampled in the natural scale.

We can now proceed with fitting, checking convergence, and plotting the posteriors and inspecting credible intervals. Note that we use the same functions as in the baseline model examples in earlier sections.

```
emc <- fit(emc)
check(emc)
```

```
#> Iterations:
#>      preburn burn adapt sample
#> [1,]      0      0      0  1000
#> [2,]      0      0      0  1000
#> [3,]      0      0      0  1000
#>
#> mu
#>      B      v    v_lMd    t0 B.w
#> Rhat  1.001  1.001  1.001    1  1
#> ESS 2361.000 2266.000 2448.000 2129 2148

#>
#> sigma2
#>      B      v    v_lMd    t0 B.w
#> Rhat  1.001  1.002  1.001  1.002  1.011
#> ESS 2074.000 1273.000 1616.000 1273.000 1070.000

#>
#> alpha highest Rhat : 3
#>      B      v    v_lMd    t0 B.w
#> Rhat  1.006  1.003  1.002  1.006  1
#> ESS 935.000 1332.000 1458.000 810.000 1521
```




Figure 11: Trace plots of the sampled parameters with highest Rhat values (i.e., worst convergence) at the group-level mean (left, μ), the group-level variance (middle, σ^2), and the subject-level (α ; in this case, subject number 3). ESS = effective sample size.

```
plot_pars(emc)
```




Figure 12: Posterior density (black) and prior density (red) of sampled group-level mean parameters of the RDM with a linear trend. Contraction indicates the reduction of uncertainty in the posterior versus the prior (1 minus the ratio of the posterior to prior variance).

```
credint(emc)
```

```
#> $mu
#>      2.5%    50%  97.5%
#> B      1.082  1.446  1.918
#> v      0.391  0.709  1.018
#> v_lMd  0.886  1.489  2.059
#> t0     -2.020 -1.701 -1.455
#> B.w    -0.065  0.098  0.256
```

On a group-level, we find a small (not credible) positive *B.w* – on average, in this dataset, thresholds appear to increase by ~ 0.098 over the course of 1024 trials, a fairly minor change relative to the intercept of ~ 1.446 . We can now generate posterior predictives, and while doing that, ask *EMC2* to return the mapped and transformed parameters, which will allow us to visualize the fitted thresholds over trials. Here we see that although there is little change for some participants, for others (e.g., 6) the change is more substantial. We can also plot the individual thresholds of each posterior predictive dataset separately as lines by setting `pp_shaded=FALSE` in the plotting function (results not shown for brevity).

```
# return_trialwise_parameters=TRUE stores the trial-by-trial parameter values as an attribute
# of the returned object, accessible via attr(pp, "trialwise_parameters")
pp <- predict(emc, return_trialwise_parameters=TRUE)
plot_trend(attr(pp, "trialwise_parameters"), emc=emc,
            par_name="B",
            filter=function(data) data$LR=="odd")
```




Figure 13: Estimated thresholds (B) as a function of trial number. Each panel corresponds to one participant (titles). Green shaded areas correspond to the 95% credible interval, the dark green line corresponds to the mean across posterior predictive datasets.

Another question we might have is whether the model with linearly changing thresholds is more supported than the baseline model. We can again use `compare()` to find out:

```
compare(list(baseline=emc_rdm, linear=emc))
```

```
#>           MD wMD   DIC wDIC  BPIC wBPIC EffectiveN meanD Dmean  minD
#> baseline -6577   0 -6680    0 -6651     0         29 -6708 -6732 -6737
#> linear   -6630   1 -6750    1 -6714     1         35 -6785 -6815 -6820
```

The results demonstrate that the complexity is warranted according to the three metrics.

Using the same logic as before, we can perform a parameter recovery by selecting a replicate dataset with its corresponding data-generating parameters, fit the model to that dataset, and compare data-generating parameters with estimated parameters:

```
postn <- 1
simulated_data <- pp[pp$postn==postn,]
true_pars <- attr(pp, "pars")[[postn]]
true_trajectories <- attr(pp, 'trialwise_parameters')[[postn]]

# combine them again in an EMC object, and fit
emc_simulated_linB <- make_emc(simulated_data, RDM_lin_B3, prior_list=prior_linB3)
emc_simulated_linB <- fit(emc_simulated_linB)

# As before, plot recovery of the parameters
recovery(emc_simulated_linB, true_pars, selection="alpha")
```




Figure 14: Parameter recovery results for the subject-level parameters. Estimated parameters (y-axis) are plotted against data-generating (x-axis). Error bars indicate 95% credible intervals, circles are on the median posterior, and diagonal lines is the identity. Pearson correlation between data-generating and posterior median. RMSE = root mean squared error.

We can also plot the recovery of the parameter trajectories by generating new posterior predictives (with estimated trajectories) from the model fit, and additionally providing the data-generating trajectories:

```
pp2 <- predict(emc_simulated_linB, return_trialwise_parameters=TRUE)
plot_trend(attr(pp2, 'trialwise_parameters'), emc=emc_simulated_linB, par_name='B',
           true_trajectories=true_trajectories)
```




Figure 15: Estimated thresholds (B) as a function of trial number. Each panel corresponds to one participant (titles). Green shaded areas correspond to the 95% credible interval, the dark green line corresponds to the mean across posterior predictive datasets. The dotted black lines are the true, data-generating trajectories of the thresholds.

0.2.2 Exponential and power laws

Practice effects tend to be large initially and then gradually decay to a positive asymptote. Exponential or power functions can be used to model such effects (Heathcote, Brown, and Mewhort 2000; Evans et al. 2018). *EMC2* offers kernels that enforce either an increase or decrease. To demonstrate these functional forms:

```

kernel_exp_incr <- make_kernel(type="exp_incr", cov_names="trials2")
kernel_exp_decr <- make_kernel(type="exp_decr", cov_names="trials2")
base_exp_incr <- make_base(type="lin", target_parameter="B", kernel=kernel_exp_incr)
base_exp_decr <- make_base(type="lin", target_parameter="B", kernel=kernel_exp_decr)
trend_exp_incr <- make_trend(base_exp_incr)
trend_exp_decr <- make_trend(base_exp_decr)

# The two designs below are identical but for the trend argument
design_exp_incr <- design(model=RDM,
  data=dat,
  contrasts=list(lM=ADmat),
  covariates="trials2",
  matchfun=function(d) d$S==d$lR,
  transform=list(func=c(B="identity")),
  formula=list(B~1, v~lM, t0~1),
  trend=trend_exp_incr,
  report_p_vector=FALSE)
design_exp_decr <- design(model=RDM,
  data=dat,
  contrasts=list(lM=ADmat),
  covariates="trials2",
  matchfun=function(d) d$S==d$lR,
  transform=list(func=c(B="identity")),
  formula=list(B~1, v~lM, t0~1),
  trend=trend_exp_decr,
  report_p_vector=FALSE)
emc_incr <- make_emc(dat, design_exp_incr)
emc_decr <- make_emc(dat, design_exp_decr)

# B.d_ei and B.d_ed are the decay parameters for the increasing and decreasing exponential
# kernels, respectively
p_vector_incr <- c(B=1, v=1, v_lMd=1, t0=0.1, B.w=1, B.d_ei=2)
p_vector_decr <- c(B=1, v=1, v_lMd=1, t0=0.1, B.w=1, B.d_ed=2)

par(mfrow=c(1, 2))
plot_trend(p_vector_incr, emc=emc_incr,
  par_name="B", subject=1, filter=function(data) data$lR=="odd",
  main="B odd")
plot_trend(p_vector_decr, emc=emc_decr,
  par_name="B", subject=1, filter=function(data) data$lR=="odd",
  main="B odd")

```




Figure 16: Example of increasing (left) and decreasing (right) exponential trends on threshold.

Note that the interpretation of the base (B) and weight ($B \cdot w$) parameters depend on the direction. In the increasing case, B corresponds to the intercept and the asymptote is $B + B \cdot w$; in the decreasing, B is the asymptote and $B + B \cdot w$ is the intercept.

0.2.2.1 Advanced note: Enforcing a direction Having both increasing and decreasing kernels may appear redundant, since the direction of the effect can also be flipped by the w parameter in the `lin` base. However, having both increasing and decreasing kernels facilitates enforcing a directional effect. For example, one might hypothesise that `v_lMd` (i.e., the ability to discriminate odd from even) increases sharply over the first few trials due to practice effects, and then stabilises. An increase could be implemented with `exp_incr`, but if the sampler converges on negative values of w of the base, the resulting trend is actually decreasing rather than increasing.

We can force the sampling of only increasing trends by restricting the range of w . This can be done by sampling w on the log-scale, and transforming it to the natural scale (see below for a demonstration). Note that transformations applied to parameters relating to the trends are always applied prior to estimating the trend.

```

trend_exp_incr <- make_trend(make_base(type="lin",
                                     target_parameter="v_lMd",
                                     kernel=kernel_exp_incr))

design_exp_incr <- design(model=RDM,
                        data=dat,
                        contrasts=list(lM=ADmat),
                        covariates="trials2",
                        matchfun=function(d) d$S==d$L,
                        # v_lMd.w="exp" samples the weight on the log scale, ensuring it is
                        always positive and thus enforcing an increasing trend regardless of
                        the sampled value
                        transform=list(func=c(v="identity",
                                              v_lMd.w="exp")),
                        formula=list(B~1, v~lM, t0~1),
                        trend=trend_exp_incr,
                        report_p_vector=FALSE)

mapped_pars(design_exp_incr)
emc_incr <- make_emc(dat, design_exp_incr)

# Three p_vectors with different values of v_lMd.w to illustrate that the direction of the
trend is always increasing
p_vector1 <- c(B=1, v=3, v_lMd=1, t0=0.1, v_lMd.w=-1, v_lMd.d_ei=2)
p_vector2 <- c(B=1, v=3, v_lMd=1, t0=0.1, v_lMd.w= 0, v_lMd.d_ei=2)
p_vector3 <- c(B=1, v=3, v_lMd=1, t0=0.1, v_lMd.w= 1, v_lMd.d_ei=2)

par(mfcol=c(2,3))
plot_trend(p_vector1, emc=emc_incr, par_name="v",
           subject=1, filter=function(data) data$lM==TRUE,
           main="v matching", ylim=c(3.5, 5))
plot_trend(p_vector1, emc=emc_incr, par_name="v",
           subject=1, filter=function(data) data$lM==FALSE,
           main="v mismatching", ylim=c(1, 2.5))
plot_trend(p_vector2, emc=emc_incr, par_name="v",
           subject=1, filter=function(data) data$lM==TRUE,
           main="v matching", ylim=c(3.5, 5))
plot_trend(p_vector2, emc=emc_incr, par_name="v",
           subject=1, filter=function(data) data$lM==FALSE,
           main="v mismatching", ylim=c(1, 2.5))
plot_trend(p_vector3, emc=emc_incr, par_name="v",
           subject=1, filter=function(data) data$lM==TRUE,
           main="v matching", ylim=c(3.5, 5))
plot_trend(p_vector3, emc=emc_incr, par_name="v",
           subject=1, filter=function(data) data$lM==FALSE,
           main="v mismatching", ylim=c(1, 2.5))

```

```

#> $v
#> lM
#> TRUE : v + 0.5 * v_lMd_t
#> FALSE : v - 0.5 * v_lMd_t
#> Trends:
#> v_lMd_t = v_lMd + w * (1 - exp(-exp(v_lMd.d_ei) * trials2))

```




Figure 17: Demonstration of enforcing a direction of an effect. Irrespective of the sign of $v_lmd.w$ (-1 for column 1, 0 for column 2, and 1 for column 3), the difference between the matching and mismatching accumulator’s drift rate increases over trials when $v_lmd.w$ is estimated on the log scale and the increasing exponential kernel is used.

In this set-up, no matter the value of $v_lmd.w$ the between-accumulator difference in drift rates increases over time. With the same logic we can also enforce a decrease by using `exp_decr` as a kernel and by applying an exponential transform to w .

0.2.3 Experimental factors on trend parameters

So far, we only inspected the effect of time on task as defined by trial number. Other options include time on task within block (e.g., short breaks between trials could cause an increased threshold for the first few trials in each block), or the number of times a stimulus digit has been shown. Furthermore, the parameters describing the trend could differ between blocks and stimulus digits. To test such hypotheses, we can allow the trend parameters themselves to vary with experimental factors, as defined in `formula` in `design()`. The example data does not have stimulus digit recorded, and all trials were run in a single block, but to provide a demonstration we simulate such effects. We then have thresholds decreasing exponentially with trial number within block, but with different rates between blocks. Additionally, we have the quality of evidence v_lmd increase exponentially with digit repetition:

```

# simulate block number (assume 3 blocks)
dat$block <- as.factor(as.numeric(cut(dat$trials, breaks=3)))
for(subject in unique(dat$subjects)) {
  dat[dat$subjects==subject, "trial_in_block"] <- ave(
    seq_along(dat[dat$subjects==subject, "block"]),
    dat[dat$subjects==subject, "block"], FUN=seq_along)/1024
}
# simulate presented digit
dat$stimulus_repetition <- ave(seq_along(dat$S),
                              dat$S, FUN=seq_along)

# combine two trends
kernel1 <- make_kernel(type="exp_decr", cov_names="trial_in_block")
kernel2 <- make_kernel(type="exp_incr", cov_names="stimulus_repetition")
base1 <- make_base(type="lin", target_parameter="B", kernel=kernel1)
base2 <- make_base(type="lin", target_parameter="v_lMd", kernel=kernel2)
trends <- make_trend(base1, base2)

design_multitrend <- design(model=RDM,
                          data=dat,
                          contrasts=list(lm=ADmat),
                          covariates=c("trial_in_block",
                                        "stimulus_repetition"),
                          matchfun=function(d) d$S==d$LR,
                          transform=list(func=c(v_lMd.w="exp",
                                                  B.w="exp",
                                                  v="identity",
                                                  B="identity")),
                          # B.d_ed~block allows the decay parameter of the threshold trend to
                          # vary between blocks (i.e., each block gets its own decay rate)
                          formula=list(B~1, v~LM, t0~1, `B.d_ed`~block),
                          trend=trends,
                          report_p_vector=FALSE)
emc_multitrend <- make_emc(dat, design_multitrend)

# thresholds
p_vector <- c(B=1, v=1, v_lMd=2, t0=0.1,
             B.d_ed=2, B.d_ed_block2=1, B.d_ed_block3=1, B.w=.25,
             v_lMd.w=1, v_lMd.d_ei=.2)

par(mfcol=c(1, 3))
plot_trend(p_vector, emc=emc_multitrend, par_name="B",
           subject=1, filter=function(data) data$LR=="odd",
           main="Thresholds", xlab="Trial number")

# on_x_axis overrides the default x-axis (trial number) with a different covariate
plot_trend(p_vector, emc=emc_multitrend, par_name="v",
           subject=1, filter=function(data) data$LM==TRUE,
           main="Drift rate (matching)", on_x_axis="stimulus_repetition",
           xlab="Stimulus repetition", xlim=c(1, 20))
plot_trend(p_vector, emc=emc_multitrend, par_name="v",
           subject=1, filter=function(data) data$LM==TRUE,
           main="Drift rate (matching)", xlab="Trial number", xlim=c(1, 20))
mapped_pars(design_multitrend)

```

```

#> $B
#>
#> intercept : B_t

```

```

#> Trends:
#> B_t = B + w * exp(-exp(B.d_ed) * trial_in_block)
#>
#> $v
#> 1M
#> TRUE : v + 0.5 * v_lMd_t
#> FALSE : v - 0.5 * v_lMd_t
#> Trends:
#> v_lMd_t = v_lMd + w * (1 - exp(-exp(v_lMd.d_ei) * stimulus_repetition))
#>
#> $B.d_ed
#> block
#> 1 : exp(B.d_ed)
#> 2 : exp(B.d_ed + B.d_ed_block2)
#> 3 : exp(B.d_ed + B.d_ed_block3)

```




Figure 18: Effect of specifying multiple trends with parameters that vary between blocks (left) or conditions (middle, right).

The left panel illustrates how the thresholds decrease within each block (at different rates between blocks), but then reset to the same asymptote after every break. The middle panel shows the drift rate for the correct accumulator as a function of how often a stimulus has been presented. Note that, when plotted against trial number (right panel), this leads to strong variability in rates across trials.

0.2.4 Polynomials

When a researcher has no strong prior assumptions about the shape of the trend but wishes to capture smooth changes, it is possible to fit a set of basis functions, such as polynomials. *EMC2* currently includes second-, third-, and fourth-order polynomials as available kernels. Many functions can be approximated with increasing accuracy by summing polynomials of increasing order. However, the increased accuracy comes at a cost of increased sampling uncertainty, since more parameters need to be estimated.

```
trend_help(kernel="poly4")
```

```

#> Description:
#> Quartic polynomial: k = d1 * c + d2 * c^2 + d3 * c^3 + d4 * c^4
#>
#> Default transformations (in order):
#> list(d1 = "identity", d2 = "identity", d3 = "identity", d4 = "identity")
#>
#> Available bases, first is the default:

```

```
#> add, identity
#>
```

Note that the kernel parameters (d1, d2, d3, d4) already weigh the contributions of the individual polynomial terms. As such, this kernel should not be combined with a linear base (`base="lin"`), but instead with an additive base (`base="add"`). Additionally, we should center the polynomial basis around the mean of the covariate (allowing there to be changes in the direction of the trend, which is not possible if the covariate is always positive). The following demonstrates the flexibility of a 4th order polynomial with a few parameter sets chosen to demonstrate the flexibility of this approach:

```
dat$tctrd <- dat$trials2-mean(dat$trials2)
# for a single base & kernel, a one-liner suffices
trend_poly <- make_trend(make_base(type="add", target_parameter="B",
                                   kernel=make_kernel(type="poly4",
                                                       cov_names="tctrd")))

design_poly <- design(model=RDM,
                     data=dat,
                     contrasts=list(lm=ADmat),
                     covariates="tctrd",
                     matchfun=function(d) d$S==d$lR,
                     transform=list(func=c(B="identity")),
                     formula=list(B~1, v~lm, t0~1),
                     trend=trend_poly,
                     report_p_vector=FALSE)
mapped_pars(design_poly)
emc_poly <- make_emc(dat, design_poly)

p_vector1 <- c(B=1, v=1, v_lMd=2, t0=0.1, B.d1= 2, B.d2= 3, B.d3=-5, B.d4= -5)
p_vector2 <- c(B=1, v=1, v_lMd=2, t0=0.1, B.d1=-1, B.d2= 0, B.d3=-5, B.d4= 0)
p_vector3 <- c(B=1, v=1, v_lMd=2, t0=0.1, B.d1= 0, B.d2= 2, B.d3= 8, B.d4= 8)
p_vector4 <- c(B=1, v=1, v_lMd=2, t0=0.1, B.d1= 1, B.d2=-3, B.d3=-2, B.d4= 10)

par(mfrow=c(1, 4))
plot_trend(p_vector1, emc=emc_poly, par_name="B",
           subject=1, filter=function(data) data$lR=="odd",
           main="p_vector1", xlab="Trial number")
plot_trend(p_vector2, emc=emc_poly, par_name="B",
           subject=1, filter=function(data) data$lR=="odd",
           main="p_vector2", xlab="Trial number")
plot_trend(p_vector3, emc=emc_poly, par_name="B",
           subject=1, filter=function(data) data$lR=="odd",
           main="p_vector3", xlab="Trial number")
plot_trend(p_vector4, emc=emc_poly, par_name="B",
           subject=1, filter=function(data) data$lR=="odd",
           main="p_vector4", xlab="Trial number")
```

```
#> $B
#>
#> intercept : B_t
#> Trends:
#> B_t = B + (B.d1 * tctrd + B.d2 * tctrd^2 + B.d3 * tctrd^3 + B.d4 * tctrd^4)
#>
#> $v
#> lm
#> TRUE : exp(v + 0.5 * v_lMd)
#> FALSE : exp(v - 0.5 * v_lMd)
```




Figure 19: Illustration of the flexibility of the polynomial kernel: With four different (arbitrarily chosen) parameter sets, widely different shapes of trends can be obtained. Here, they are visualised as trends on thresholds B , but they can be on any parameter.

0.3 Mechanisms of dynamics

Models with dynamic mechanisms (as opposed to the descriptions of dynamics in the last section) often rely on delta rules to learn quantities (here referred to generically as Q) used for trial-to-trial control of some aspect of the decision process. The implementation follows the same general logic as above — the delta rule is implemented as a kernel in `make_kernel()`. Below, we provide four demonstrations of such mechanisms spanning different learnt capacities and their effects on parameters.

0.3.1 Stimulus memory

To illustrate we implement the stimulus memory mechanism proposed by Miletic et al. (2025). Here, the decision maker keeps track of the probability that a stimulus is even (or odd). The covariate in this case is the stimulus, which we recode as 1 (even) and -1 (odd) and use in the delta rule:

$$Q_{SM,t+1} = Q_{SM,t} + \alpha_{SM} \cdot (S_t - Q_{SM,t}), \quad S \in \{-1, 1\}$$

The delta rule introduces two extra parameters: a learning rate α_{SM} , and a value for Q at the start of the experiment $Q_{SM,0}$. The latter is hard to estimate and often has little influence (as long as the learning rate is not very low), so we tend to fix it to a constant—for this rule, a constant of 0 implies a belief that both stimuli occur equally often.

For now, let us assume that stimulus memory controls response bias as mediated by the relative thresholds, so an increase in the threshold for the ‘even’ accumulator is accompanied by an equal and opposite decrease for the ‘odd’ accumulator (and vice versa). To achieve this, we need to parametrise thresholds with the same mean-difference parametrisation as we used for drift rates. That is,

$$\begin{aligned} B_{odd} &= B_{mean} + B_{lRd} \\ B_{even} &= B_{mean} - B_{lRd} \end{aligned}$$

The B_{lRd} term is the parameter influenced by the updated covariate — however, since we do not necessarily expect an across-trial *mean* difference between thresholds, we set B_{lRd} as a constant to 0.

Using a linear base, we then allow the threshold difference to vary on a trial-by-trial basis with

$$B_{lRd,t} = B_{lRd} + w_{SM} \cdot Q_{SM,t}$$

Note the extra parameter here: w_{SM} .

```
# Let's reload the data to drop all covariates created earlier
load(url(data_url))
dat$Stim1 <- ifelse(dat$S=="even", 1, -1)
SM_kernel <- make_kernel(type="delta", cov_names="Stim1")
SM_base <- make_base(type="lin", target_parameter="B_lRd",
                     kernel=SM_kernel, phase="premap")
SM_trend <- make_trend(SM_base)

RDM_SM <- design(model=RDM,
                 data=dat,
                 contrasts=list(lM=ADmat, lR=ADmat),
                 covariates="Stim1",
                 matchfun=function(d) d$S==d$lR,
                 formula=list(B~lR, v~lM, t0~1),
                 trend=SM_trend,
                 constants=c(B_lRd=0, B_lRd.q0=0),
                 report_p_vector=FALSE)
```

We can now apply a similar prior to B as above and fit the model. Note that running `make_emc()` in the code snippet below returns a message saying compression is automatically turned off. Compression is a default feature in *EMC2* (Stevenson et al. 2026) which leverages redundancy in choice-RT data to speed up likelihood computations. However, compression is not possible with delta rules, which require all trials, and so it is turned off once a delta rule trend is detected.

```
prior_SMB <- prior(RDM_SM, mu_mean=c(B=2), mu_sd=c(B=0.5))
emc <- make_emc(dat, design=RDM_SM, prior_list=prior_SMB)
emc <- fit(emc)
check(emc)
```

```
#> Iterations:
#>      preburn burn adapt sample
#> [1,]      0      0      0    1000
#> [2,]      0      0      0    1000
#> [3,]      0      0      0    1000
#>
#> mu
#>      B      v v_lMd      t0 B_lRd.alpha B_lRd.w
#> Rhat  1.019  1.001    1  1.004      1.006  1.006
#> ESS 1065.000 1464.000 1773 1694.000    869.000 1429.000

#>
#> sigma2
#>      B      v      v_lMd      t0 B_lRd.alpha B_lRd.w
#> Rhat  1.008  1.003  1.001  1.002      1.038  1.016
#> ESS  923.000 1207.000 1421.000 1248.000    415.000 743.000

#>
#> alpha highest Rhat : 5
#>      B      v      v_lMd      t0 B_lRd.alpha B_lRd.w
#> Rhat  1.002  1.005  1.004  1.002      1.02  1.018
#> ESS 1811.000 1035.000 950.000 1840.000    304.00 408.000
```




Figure 20: Trace plots of the sampled parameters with highest Rhat values (i.e., worst convergence) at the group-level mean (left, μ), the group-level variance (middle, σ^2), and the subject-level (α ; in this case, subject number 5). ESS = effective sample size.

```
plot_pars(emc)
```




Figure 21: Posterior density (black) and prior density (red) of sampled group-level mean parameters of the RDM with an SM mechanism. Contraction indicates the reduction of uncertainty in the posterior versus the prior (1 minus the ratio of the posterior to prior variance).

We find excellent convergence properties, especially given that learning rate parameters are often hard to estimate. Note that learning rates are by default sampled on the probit scale, so if we want to know the mean learning rate, we need to map and transform the parameters first (this can take a while to run):

```
credint(emc, map=TRUE)
```

```
#> $mu
#>          2.5%  50% 97.5%
#> v_lMFALSE 0.839 0.987 1.144
#> v_lMTRUE  4.250 4.365 4.483
#> B_lReven  1.291 1.349 1.410
#> B_lRodd   1.321 1.379 1.441
#> t0        0.208 0.215 0.222
```

```
#> B_1Rd.w      0.276 0.322 0.374
#> B_1Rd.alpha 0.321 0.388 0.462
```

Hence, we find a group-wise mean learning rate of ~ 0.388 , and the threshold of the odd accumulator is 0.322 lower than the threshold for the even accumulator if the participant has a Q_{SM} -value of 1.

We can now check whether the SM mechanism can explain stimulus-history effects by plotting the RTs (by choice) and the probability of choosing ‘even’ as a function of the previous three presented stimuli. First, generate posterior predictives, and then plot the stimulus history effects:

```
pp <- predict(emc, return_trialwise_parameters=TRUE)
# this is a convenience function that is not part of EMC2. For help, inspect the first lines
# after calling plot_history_effects (omitting the parentheses), which document its use
plot_history_effects(dat, pp, n_hist=3)
```




Figure 22: Effects of local stimulus (S) history on response times (left) and response (R) probability (right). X-axis labels indicate the three most recent trials’ stimuli (e.g., eee = even, even even, eeo = even, even, odd; the right-most letter in the string indicates the most recent trial). Crosses indicate data, model predictions are indicated by error bars (95% credible intervals), and lines connect medians. Red circles in the right plot indicate the probability of the stimulus, which is an important quality check.

The left panel shows the mean RT as a function of the stimuli of the three previous trials (e.g., ooo = odd, odd, odd; ooe = odd, odd, even; etc, where the right-most letter in the string indicates the most recent trial). The RT data show a cross-over pattern, such that if the previous three stimuli were odd, ‘odd’ responses are faster than ‘even’ responses; and vice versa. The model captures this cross-over pattern fairly well, with some misfits mostly due to an apparent asymmetry in the cross-over in the data.

The right panel shows the probability of choosing even (black) and the probability of the stimulus being even (red) as a function of local trial history. Specifically, for each trial history triplet (e.g., eee, eeo, etc.), the red circles reflect the proportion of trials in which the current stimulus was ‘even’, computed across all trials sharing that three-trial history — that is, the empirical conditional probability $P(\text{stimulus}=\text{even}|\text{trial history})$. The black crosses demonstrate that participants are more likely ($\approx .55$) to choose ‘even’ if the previous three trials were ‘even’. The red circles are important as a quality check. If the stimuli are generated truly randomly (as is the case in this experiment), the red circles should be approximately 0.5 in all cases, since the probability of the stimulus being even was intended to be 0.5 in this experiment. However, some experiments use pseudo-randomisation, which can lead to negative correlations in local stimulus histories. This can confound any stimulus history effects present in the human data. True stimulus history effects are those that are larger than the ones present in the data, e.g., visible as black crosses deviating more from 0.5 compared to red circles.

As before, we can plot the trial-wise parameters with `plot_trend()`. Let’s zoom in on the even and odd thresholds of the first three subjects for the first 20 trials:

```

par(mfcol=c(2,3))
for(subject in 1:3) {
  plot_trend(attr(pp, "trialwise_parameters"), emc=emc,
    par_name="B", subject=subject,
    filter= function(data) data$lR=="odd",
    main=paste0(subject, " odd", xlim=c(1, 20))
  plot_trend(attr(pp, "trialwise_parameters"), emc=emc,
    par_name="B", subject=subject,
    filter= function(data) data$lR=="even",
    main=paste0(subject, " even", xlim=c(1, 20))
}

```




Figure 23: Estimated trialwise thresholds of the accumulator corresponding to responding odd (top) and even (bottom) for the first three participants (columns). Green shaded areas indicate the 95% credible interval.

Alternatively, we can plot the evolution of the Q-values themselves. Note that `predict()` was run with another argument: `return_trialwise_parameters=TRUE`, which returns the updated covariates as an attribute. We can use this as follows:

```

# k1.Qvalue is the name EMC2 assigns to the Q-values of the delta rule of kernel 1
plot_trend(attr(pp, "trialwise_parameters"), emc=emc,
  par_name="k1.Stim1.Qvalue",
  xlim=c(1, 20))

```




Figure 24: Estimated Q-values corresponding to the SM mechanism for all six participants (panels). Green shaded areas indicate the 95% credible interval.

0.3.1.1 Learning in multiple-accumulator models Under the hood, EMC2 works with data-augmented design matrices (`dadms`). A `dadm` contains one row per accumulator per trial – so for race models applied to two-alternative forced choice tasks, this will typically result in two rows per trial. We can inspect the `dadm` for the first subject in the `emc` object to see that the covariate `Stim1` is indeed the same across accumulators in each trial:

```
head(emc[[1]]$data[[1]])
```

```
#>  subjects    S    R   rt Stim1 trials   lR   lM winner
#> 1      1 even even 0.542    1     1 even  TRUE   TRUE
#> 2      1 even even 0.542    1     1 odd  FALSE  FALSE
#> 3      1 even even 0.426    1     2 even  TRUE   TRUE
#> 4      1 even even 0.426    1     2 odd  FALSE  FALSE
#> 5      1 even even 0.501    1     3 even  TRUE   TRUE
#> 6      1 even even 0.501    1     3 odd  FALSE  FALSE
```

The delta rule is applied to the covariate as it is specified in the `dadm`. Since, in the examples in this tutorial, the covariate is the same for both accumulators, applying a delta rule to the covariate in the `dadm` would lead to *two* updates every trial: One for the first accumulator, and then another one for the second accumulator. The Q-values would then also differ between accumulators. The correct behavior is to update only once every trial (i.e., for the first accumulator), and then feed forward the updated Q-value to every second (and third, fourth, ...) accumulator. By setting `at="lR"`, the Q-value is only updated at the first level of the factor in `lR`, and the resulting Q-value is carried over to the other levels of `lR`.

In many applications when using dynamic EAMs, this is likely the intended behavior, and therefore `lR` is the default setting for `at`. The DDM marks an exception, since it assumes one accumulator per trial, so has no `lR` factor, as we demonstrate below. Advanced use cases might need to set `at=NULL`.

0.3.2 Repetition/alternation memory

The stimulus-memory mechanism accounts for what Jones et al. (2013) termed *first-degree* recency effects: They depend on stimulus identity (e.g., ‘odd’ or ‘even’). There are also *second-degree* recency effects, which depend on sequential properties - i.e., whether stimulus identities are repetitions or alternations of those

stimuli before. With an appropriate representation it is relatively easy to implement such a mechanism with a delta rule as well. If we code the stimuli as $S_t \in -1, 1$, then the product $S_{t-1}S_t$ indicates if trial t repeats (outcome 1) or alternates (outcome -1) trial $t - 1$. Let's hypothesise that participants keep track of the repetition rate using a repetition memory (RM) based on the delta rule:

$$Q_{RM,t+1} = Q_{RM,t} + \alpha_{RM} \cdot (S_{t-1}S_t - Q_{RM,t}), \quad S \in \{-1, 1\}$$

Q_{RM} is a latent representation for the belief that the stimulus will repeat itself. We could hypothesise that, if participants believe that the stimulus is repeating, the threshold corresponding to the choice option that repeats the previous stimulus is lower compared to the one that alternates - and vice versa. To implement this, we need to parametrise the thresholds as repeat/alternate:

$$\begin{aligned} B_{repeat} &= B_{mean} + B_{lR2d} \\ B_{alternate} &= B_{mean} - B_{lR2d} \end{aligned}$$

As before, we are making use of a mean-difference parametrisation for the thresholds here, but the difference B_{lR2d} parameter here codes accumulators by whether they repeat or alternate the previous stimuli. This B_{lR2d} term is the parameter influenced by the updated covariate. Since we do not necessarily expect an across-trial *mean* difference between thresholds, we set B_{lR2d} itself as a constant to 0.

Using a linear base, we then allow the threshold difference to vary on a trial-by-trial basis with

$$B_{lR2d,t} = B_{lR2d} + w_{RM} \cdot Q_{RM,t}$$

Again, this introduces a weighting parameter: w_{RM} . Let's implement this step by step. First, we add $S_{t-1}S_t$ to the data, as `dat$Stim2`. This will be our new covariate. It should affect `B_1R2d` with a linear basis. Unlike `1R`, the `1R2` factor is not automatically created by *EMC2*, so we must pass a function to design that creates it for us.

```

# function to lag
lag <- function(x, shift = 1) {
  n <- length(x)
  if (shift==0) return(x)

  ret <- rep(NA, n)
  if (abs(shift) < n) {
    if (shift > 0) {
      ret[-seq_len(shift)] <- x[seq_len(n-shift)]
    } else {
      ret[seq_len(n+shift)] <- x[(1-shift):n]
    }
  }
}

attributes(ret) <- attributes(x)
ret
}

dat$Stim1 <- ifelse(dat$S=="even", 1, -1)      # S_{t}
dat$Stim2 <- dat$Stim1*lag(dat$Stim1,1) # S_{t-1}S_{t}
RM_kernel <- make_kernel(type="delta", cov_names="Stim2")
RM_base <- make_base(type="lin", target_parameter="B_lR2d",
                     kernel=RM_kernel, phase="premap")
RM_trend <- make_trend(RM_base)

make_lR2 <- function(x) {
  # note that we need to use lag 2, since the dadm has one row per accumulator
  lR2 <- as.numeric(x$lR)==as.numeric(lag(x$S, 2))
  # The first trial has no repeat/alternate. However, since we set B_lR2d.q0, we can set lR2
  # for this trial arbitrarily TRUE or FALSE -- B_lR2d_{1} is always 0.
  lR2[is.na(lR2)] <- TRUE
  factor(lR2, levels=c(1, 0), labels=c("rep", "alt"))
}

RDM_RM <- design(model=RDM,
                 data=dat,
                 contrasts=list(lM=ADmat, lR=ADmat, lR2=ADmat),
                 covariates="Stim2",
                 functions=list(lR2=make_lR2),
                 matchfun=function(d) d$S==d$lR,
                 formula=list(B~lR2, v~lM, t0~1),
                 trend=RM_trend,
                 constants=c(B_lR2d=0, B_lR2d.q0=0),
                 report_p_vector=FALSE)

```

Before fitting, we might want to simulate some data to test whether it indeed demonstrates second-degree recency effects.

```

p_vector <- sampled_pars(RDM_RM)
p_vector[1:6] <- c(log(2), log(1), .5, log(.15), qnorm(0.1), 2)
simDat <- make_data(p_vector, RDM_RM, data=dat)
# the plot_history_effects()-function takes an additional argument: degree, which can be 1 for
# first-degree or 2 for second-degree effects
plot_history_effects(dat=simDat, degree=2)

```




Figure 25: Effects of local stimulus (S) repetition history on response times (left) and response (R) probability (right). X-axis labels indicate the three most recent trials' stimuli, coded as whether they repeated (R) or alternated (A) the preceding trial (e.g., AAA = alternate, alternate, alternate, AAR = alternate, alternate, repeat; the right-most letter in the string indicates the most recent trial). Crosses indicate data. Red circles in the right plot indicate the probability of the stimulus repeating the previous stimulus, which is an important quality check.

This shows the expected pattern: if the previous trials were repeats, the model predicts participants are more likely to repeat on the current trial. Furthermore, a “repeat” response would be faster than an “alternate” response. Now let's fit the model to see the strength of these effects in real data:

```
emc <- make_emc(dat, design=RDM_RM)
emc <- fit(emc)
pp <- predict(emc)
plot_history_effects(dat, pp, degree=2)
```




Figure 26: Effects of local stimulus (S) repetition history on response times (left) and response (R) probability (right). X-axis labels indicate the three most recent trials' stimuli, coded as whether they repeated (R) or alternated (A) the preceding trial (e.g., AAA = alternate, alternate, alternate, AAR = alternate, alternate, repeat; the right-most letter in the string indicates the most recent trial). Crosses indicate data, model predictions are indicated by error bars (95% credible intervals), and lines connect medians. Red circles in the right plot indicate the probability of the stimulus repeating the previous stimulus, which is an important quality check.

These results demonstrate that the second-degree recency effects in response probability are fit reasonably well, but some misfit remains in the response times.

0.3.2.1 WDM and drift rate effects The examples above follow Miletić et al. (2025) by using an RDM and assuming that stimulus memory causes a threshold (start point) bias. *EMC2* is flexibly designed to implement alternative hypotheses. For example, we can propose a WDM instead, and hypothesise that drift rates rather than start points are biased. However, as the minimal implementation below demonstrates, this does not appear to fit as well, and loses quantitative model comparisons:

```
SM_kernel_DDM <- make_kernel(type="delta", cov_names="Stim1", at=NULL)
SM_base_DDM <- make_base(type="lin", target_parameter="v",
                        kernel=SM_kernel_DDM, phase="premap")
SM_trend_DDM <- make_trend(SM_base_DDM)

SM_DDM <- design(model=DDM,
                data=dat,
                covariates=c("Stim1"),
                contrasts=list(S=Smat),
                formula=list(v~S, a~1, t0~1),
                constants=c(v.q0=0, v=0),
                trend=SM_trend_DDM)

# since v is estimated on the natural scale (in case of the DDM), we do not need to change the
# transformations, and we can use the default priors.
emc_ddm <- make_emc(dat, design=SM_DDM)
emc_ddm <- fit(emc_ddm)
pp_ddm <- predict(emc)

# check quality of fit
plot_history_effects(dat, pp_ddm)

# compare with RDM
compare(sList=list(rdm=emc,
                  ddm=emc_ddm))
```

```
#>
#> Sampled Parameters:
#> [1] "v_Sdif" "a" "t0" "v.alpha" "v.w"
#>
#> Design Matrices:
#> $v
#> S v v_Sdif
#> even 1 -1
#> odd 1 1
#>
#> $a
#> a
#> 1
#>
#> $t0
#> t0
#> 1
#>
#> $v.q0
#> v.q0
#> 1
#>
#> $v.alpha
#> v.alpha
#> 1
#>
```

```

#> $v.w
#> v.w
#> 1
#>
#> $s
#> s
#> 1
#>
#> $st0
#> st0
#> 1
#>
#> $sv
#> sv
#> 1
#>
#> $SZ
#> SZ
#> 1
#>
#> $Z
#> Z
#> 1

```

```

#>      MD wMD   DIC wDIC  BPIC wBPIC EffectiveN meanD Dmean minD
#> rdm -6743   1 -6862    1 -6826     1         36 -6899 -6928 -6935
#> ddm -5290   0 -5391    0 -5361     0         30 -5420 -5443 -5450

```




Figure 27: Effects of local stimulus (S) history on response times (left) and response (R) probability (right). X-axis labels indicate the three most recent trials' stimuli (e.g., eee = even, even even, eeo = even, even, odd; the right-most letter in the string indicates the most recent trial). Crosses indicate data, model predictions (in this case, the DDM with a drift rate effect of stimulus memory) are indicated by error bars (95% credible intervals), and lines connect medians. Red circles in the right plot indicate the probability of the stimulus, which is an important quality check.

0.3.3 Accuracy memory

Miletić et al. (2025) proposed an accuracy-memory mechanism that follows the same logic as above, with a different covariate (errors) and a different parameter (urgency, v). Errors are coded as the inverse of accuracy; they are 0 when the response matches the stimulus, and 1 otherwise. Unlike the stimulus frequency, accuracy/error is not a static property of the experimental design but depends on the observed behaviour. This is important when simulating new data - the covariate in question is not the *empirically-observed* accuracy on any given trial, but the *newly-simulated* accuracy.

For this reason, it is useful to have *EMC2* calculate accuracy with a function rather than specifying it in the dataframe (in which case it would be treated as a static experimental factor). The function is applied when generating posterior predictives, ensuring that the accuracy of the *simulated* data is used as a covariate.

```
AM_kernel <- make_kernel(type="delta", cov_names="error")
AM_base <- make_base(type="lin", target_parameter="v",
                    kernel=AM_kernel, phase="premap")
AM_trend <- make_trend(AM_base)

RDM_AM <- design(model=RDM,
                data=dat,
                contrasts=list(lM=ADmat, lR=ADmat),
                # functions= can be used to define covariates that depend on observed
                # behaviour (here: errors); EMC2 will recompute these from simulated data when
                # generating posterior predictives
                functions=list(error=function(x) as.numeric(x$S!=x$R)),
                matchfun=function(d) d$S==d$lR,
                # Here, we ensure that the minimum value of alpha after transformation is 0.01
                # instead of 0, to avoid regions of the parameter space where the model is
                # unidentified (with a learning rate of 0, the v.w parameter has no influence)
                # on the data
                transform=list(lower=c(v.alpha=0.01)),
                formula=list(B~1, v~lM, t0~1, s~lM),
                trend=AM_trend,
                constants=c(v.q0=0, s=1),
                report_p_vector=FALSE)

# NB1: we allow within-trial noise s to vary with accumulator match (correct/incorrect), which
# will allow us to capture the relative speed of errors.

# NB2: in this case, we are not estimating v on the natural scale:
mapped_pars(RDM_AM)
# This is because the AM mechanism can push drift rates on individual trials to negative values
# when applied on the natural scale. Since the RDM has no known analytic likelihood for negative
# drift rates these are excluded by the RDM using its bound attribute, but this can lead to
# difficulties sampling. Hence, the effect of errors on urgency as we implement it here, is not
# strictly linear -- it is linear on the log-scale. The cognitive interpretation remains similar.

emc <- make_emc(dat, design=RDM_AM)
emc <- fit(emc)
credint(emc)
```

```
#> $v
#> lM
#> TRUE : exp(v_t + 0.5 * v_lMd)
#> FALSE : exp(v_t - 0.5 * v_lMd)
#> Trends:
#> v_t = v + w * q[i]
#>
#> $s
#> lM
#> TRUE : exp(s + 0.5 * s_lMd)
#> FALSE : exp(s - 0.5 * s_lMd)

#> $mu
#> 2.5% 50% 97.5%
#> B 0.962 1.272 1.510
#> v 0.621 1.113 1.712
#> v_lMd 1.262 2.537 3.510
```

```
#> t0      -2.008 -1.724 -1.401
#> s_lMd   -0.513 -0.354 -0.155
#> v.alpha  0.126  0.780  1.360
#> v.w     -0.251 -0.163 -0.052
```

Accuracy memory can explain error-related effects such as post-error slowing. To visualise, we can generate posterior predictives and then use a convenience function to plot mean RT as a function of trial number relative to an error. As mentioned above, since the covariate in this case depends on observed behaviour, we need to simulate data and determine the covariate one trial at a time. *EMC2* detects whether the covariate is `rt`, `R`, or the output of one of the functions provided to design. If it is, it automatically simulates data trial-by-trial. Because vectorisation is not possible in this case, this way of simulating data is much slower.

```
pp <- predict(emc)
# this is a convenience function that is not part of EMC2. For help, inspect the first lines
# after calling plot_error_effects (omitting the parentheses), which document its use
plot_error_effects(dat, pp)
```




Figure 28: Quality of fit of error-related effects by an RDM that allows urgency to vary with accuracy memory. x-axes depict trial number relative to the error (trial 0); crosses are (group-averaged) mean RTs, error bars are 95% credible intervals of the models. Horizontal dashed lines indicate the mean RT.

As in the case of stimulus memory, the accuracy memory mechanism can be applied to other EAMs like the DDM, and it can affect other parameters. The DDM does not have an urgency mechanism, but we could, for example, allow accuracy memory to affect thresholds. This trend can be specified as post-transform since the thresholds are the same across all design cells. In this case, we can sample thresholds on the log scale, transform to the natural scale, and then apply the trend. However, this again does not appear to fit as well qualitatively, and loses model comparisons, as the following code shows:

```

AM_kernel_DDM <- make_kernel(type="delta", cov_names="error", at=NULL)
AM_base_DDM <- make_base(type="lin", target_parameter="a",
                          kernel=AM_kernel_DDM, phase="posttransform")
AM_trend_DDM <- make_trend(AM_base_DDM)

DDM_AM <- design(model=DDM,
                 data=dat,
                 functions=list(error=function(x) as.numeric(x$S!=x$R)),
                 contrasts=list(S=Smat),
                 formula=list(v~S, a~1, t0~1),
                 constants=c(a.q0=0),
                 trend=AM_trend_DDM,
                 report_p_vector=FALSE)

emc_ddm_am <- make_emc(dat, design=DDM_AM)
emc_ddm_am <- fit(emc_ddm_am)
pp_ddm_am <- predict(emc_ddm_am)
plot_error_effects(dat, pp_ddm_am)
compare(list(rdm=emc, ddm=emc_ddm_am))

```

```

#>      MD wMD   DIC wDIC  BPIC wBPIC EffectiveN meanD Dmean  minD
#> rdm -6903   1 -7049    1 -7012     1         37 -7086 -7112 -7123
#> ddm -5329   0 -5463    0 -5424     0         39 -5502 -5534 -5541

```




Figure 29: Quality of fit of error-related effects by a WDM that allows thresholds to vary with accuracy memory. x-axes depict trial number relative to the error (trial 0); crosses are (group-averaged) mean RTs, error bars are 95% credible intervals of the models. Horizontal dashed lines indicate the mean RT.

0.3.4 Fluency memory

Fluency memory is a formalisation of the idea that participants adjust their behaviour according to perceived difficulty (Miletić et al. 2025). To estimate perceived difficulty, we rely on a few simplifying assumptions: participants have an implicit sense of (1) their thresholds, and (2) their response time. In race models, the ratio of threshold over response time approximates mean speed, which can be used as a proxy for difficulty, as processing efficiency will be lower for difficult trials relative to easy trials. Formally, $Q_{FM,t+1} = Q_{FM,t} + \alpha_{FM}(b_t/r_t - Q_{FM,t})$.

Here, we implement this using a custom kernel. Custom kernels are fully implemented by the user in *Rcpp* (Eddelbuettel and François 2011) - here, we write the function in a separate file (path is arbitrary, but here `./custom_kernel.cpp`) with the following contents:

```

// [[Rcpp::depends(EMC2)]]
#include <Rcpp.h>
#include "EMC2/userfun.hpp"

// Example: two params (q0, alpha) and three inputs (rt, B, weight)
Rcpp::NumericVector custom_kernel(Rcpp::NumericMatrix kernel_pars,
                                  Rcpp::NumericMatrix input) {
  // assume kernels_pars(q0, alpha) and inputs(rt, B, weight)
  int n = input.nrow();
  Rcpp::NumericVector q(n);
  Rcpp::NumericVector B_t(n);
  Rcpp::NumericVector pe(n);
  // B_t = B + w*Q_t for trial 1
  q[0] = kernel_pars(0,0);
  B_t[0] = input(0, 1) + input(0, 2)*q[0];
  for (int i = 1; i < n; ++i) {
    // 'prediction error' = B_t/rt_t - Q_t
    pe[i-1] = (B_t[i-1] / input(i-1,0)) - q[i-1];
    // update: Q_t = Q_{t-1} + alpha_{t-1} * pe_{t-1}
    q[i] = q[i-1] + kernel_pars(i-1,1) * pe[i-1];
    // Apply base to B to get the correct threshold for
    // next trial (needed for PE calculation)
    B_t[i] = input(i,1) + input(i,2) * q[i];
  }
  return q;
}

// Export pointer maker for registration
// [[Rcpp::export]]
SEXP EMC2_make_custom_kernel_ptr();
EMC2_MAKE_PTR(custom_kernel);

```

We can then compile it in *R*, and supply a C pointer to the compiled kernel. Note that when we load a custom-kernel emc object from disk we need to reinitialise C pointers (see `?fix_custom_kernel_pointers`).

```

ck <- register_kernel(
  # parameter names in the order in which they are used in the C++ function
  kernel_parameters=c("q0", "alpha"),
  file="./custom_kernel.cpp",
  # alpha is a learning rate: bounded (0,1)
  transforms=c(q0="identity", alpha="pnorm")
)

FM_kernel <- make_kernel(type="custom", custom_kernel=ck,
  cov_names="rt",
  # par_input passes additional model parameters to the kernel (here: B
  # and B.w are needed to compute B_t/rt_t)
  par_input=c("B", "B.w"))
FM_base <- make_base(target_parameter="B", type="lin", kernel=FM_kernel)
trend_FM <- make_trend(FM_base)

design_FM <- design(model=RDM,
  data=dat,
  covariates=c("rt"),
  contrasts=list(lM=ADmat),
  matchfun=function(d) d$S==d$lR,
  formula=list(B~1, v~lM, t0~1, s~lM),
  transform=list(func=c(B="identity"),
    lower=c(B.alpha=0.05)),
  constants=c(B.q0=3, # NB1
    s=log(1)),
  trend=trend_FM)
# NB1: Why q0=3? When fitting a static RDM, the mean fit threshold over mean RT ratio is very
# roughly 3 in this dataset. It serves as a reasonable place to start
prior_FM <- prior(design_FM, mu_mean=c(B=2), mu_sd=c(B=0.5))
emc <- make_emc(dat, design=design_FM, prior_list=prior_FM, compress=FALSE)
emc <- fit(emc)

```

```

#>
#> Sampled Parameters:
#> [1] "B"      "v"      "v_lMd"  "t0"     "s_lMd"  "B.alpha" "B.w"
#>
#> Design Matrices:
#> $B
#> B
#> 1
#>
#> $v
#>      lM v v_lMd
#> TRUE 1  0.5
#> FALSE 1 -0.5
#>
#> $t0
#> t0
#> 1
#>
#> $s
#>      lM s s_lMd
#> TRUE 1  0.5
#> FALSE 1 -0.5
#>
#> $B.q0
#> B.q0

```

```

#>      1
#>
#> $B.alpha
#> B.alpha
#>      1
#>
#> $B.w
#> B.w
#>      1
#>
#> $A
#> A
#> 1

```

Fluency memory provides an explanation for the slower fluctuations in the observed RTs that are not addressed by the previously discussed mechanisms. One useful way to visualise these is by creating a power spectrum of the RT data, as done below. The height of the spectrum (spectral “power”) indicates the predominance of fluctuations with that period (or its inverse, frequency). The black line is the spectrum of the empirical data, and green indicates the model fit. As is commonly observed, the spectrum shows that power increases for longer periods, and that fluency can account for much of the power in fluctuations that are moderate to fast, but not the slowest fluctuations. Note that we supplied the mean trial duration to determine the periods, which are then plotted on the x axis. If this is not supplied the x axis is instead log(frequency).

```

pp <- predict(emc)
# convenience function to plot the spectrum
plot_spectrum(dat=dat, pp=pp, trial_duration=1.265)
# mean trial duration of this task was 1.265 s, which we use to determine the periods

```




Figure 30: Group-averaged power spectrum of the RT data (black line) and model posterior predictives (green). Each individual green line is one posterior predictives’ (group-averaged) power spectrum; the dark green line averages across posterior predictives.

- Ando, T. 2007. “Bayesian Predictive Information Criterion for the Evaluation of Hierarchical Bayesian and Empirical Bayes Models.” *Biometrika* 94 (2): 443–58. <https://doi.org/10.1093/biomet/asm017>.
- Eddelbuettel, Dirk, and Romain François. 2011. “Rcpp: Seamless R and C++ Integration.” *Journal of Statistical Software* 40 (8): 1–18. <https://doi.org/10.18637/jss.v040.i08>.
- Evans, Nathan J., Scott D. Brown, Douglas J. K. Mewhort, and Andrew Heathcote. 2018. “Refining the Law of Practice.” *Psychological Review* 125 (4): 592–605. <https://doi.org/10.1037/rev0000105>.
- Heathcote, Andrew, S Brown, and D J K Mewhort. 2000. “The power law repealed: The case for an exponential law of practice.” *Psychonomic Bulletin & Review* 7 (2): 185–207. <http://www.springerlink.com/index/44Q1498351121510.pdf>.

- Huang, Alan, and M. P. Wand. 2013. “Simple Marginally Noninformative Prior Distributions for Covariance Matrices.” *Bayesian Analysis* 8 (2): 439–52. <https://doi.org/10.1214/13-BA815>.
- Jones, Matt, Tim Curran, Michael C. Mozer, and Matthew H. Wilder. 2013. “Sequential Effects in Response Time Reveal Learning Mechanisms and Event Representations.” *Psychological Review* 120 (3): 628–66. <https://doi.org/10.1037/a0033180>.
- Miletić, Steven, Niek Stevenson, Ami Eidels, Dora Matzke, Birte Forstmann, and Andrew Heathcote. 2025. “Explaining Multi-Scale Choice Dynamics.” *Psychological Review*. <https://doi.org/10.1037/rev0000581>.
- Spiegelhalter, David J., Nicola G. Best, Bradley P. Carlin, and Angelika van der Linde. 2002. “Bayesian Measures of Model Complexity and Fit.” *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 64 (4): 583–639. <https://doi.org/10.1111/1467-9868.00353>.
- Stevenson, Niek, Michelle C. Donzallaz, Reilly J. Innes, Birte U. Forstmann, Dora Matzke, and Andrew Heathcote. 2026. “Bayesian Hierarchical Cognitive Modeling with the EMC2 Package.” *Behavior Research Methods* 58 (1): 35. <https://doi.org/10.3758/s13428-025-02869-y>.
- Strickland, Luke, Russell J. Boag, Niek Stevenson, and Andrew Heathcote. 2026. “An Illustrative Guide to Expressing Cognitive Theories Using Evidence Accumulation Modelling.” *Behavior Research Methods* 58 (4): 101. <https://doi.org/10.3758/s13428-026-02970-w>.
- Wagenmakers, Eric-Jan, Simon Farrell, and Roger Ratcliff. 2004. “Estimation and Interpretation of $1/\alpha$ Noise in Human Cognition.” *Psychonomic Bulletin & Review* 11 (4): 579–615. <https://doi.org/10.3758/BF03196615>.